

DB

Chapter 15: Concurrency Control

Dr. CHEN Jian
Professor
ellachen@scut.edu.cn

SCUT

15.1 Lock-Based Protocols

- A lock is a mechanism to control concurrent access to a data item
- Data items can be locked in two modes :
 1. *exclusive (X) mode*. Data item can be both read as well as written. X-lock is requested using **lock-X** instruction.
 2. *shared (S) mode*. Data item can only be read. S-lock is requested using **lock-S** instruction.
- Lock requests are made to concurrency-control manager. Transaction can proceed only after request is granted.

Lock-Based Protocols (Cont.)

■ Lock-compatibility matrix

	S	X
S	true	false
X	false	false

- A transaction may be granted a lock on an item if the requested lock is compatible with locks already held on the item by other transactions
- Any number of transactions can hold shared locks on an item,
 - but if any transaction holds an exclusive on the item no other transaction may hold any lock on the item.
- If a lock cannot be granted, the requesting transaction is made to wait till all incompatible locks held by other transactions have been released. The lock is then granted.

Lock-Based Protocols (Cont.)

- Example of a transaction performing locking:

```
 $T_2$ : lock-S(A);  
      read (A);  
      unlock(A);  
      lock-S(B);  
      read (B);  
      unlock(B);  
      display(A+B)
```

- Locking as above is not sufficient to guarantee serializability — if A and B get updated in-between the read of A and B , the displayed sum would be wrong.
- A **locking protocol** is a set of rules followed by all transactions while requesting and releasing locks. Locking protocols restrict the set of possible schedules.

Pitfalls of Lock-Based Protocols

- Consider the partial schedule

T_3	T_4
lock-X(B) read(B) $B := B - 50$ write(B) lock-X(A)	lock-S(A) read(A) lock-S(B)

- Neither T_3 nor T_4 can make progress — executing **lock-S(B)** causes T_4 to wait for T_3 to release its lock on B , while executing **lock-X(A)** causes T_3 to wait for T_4 to release its lock on A .
- Such a situation is called a **deadlock**.
 - To handle a deadlock one of T_3 or T_4 must be rolled back and its locks released.

Pitfalls of Lock-Based Protocols (Cont.)

- The potential for deadlock exists in most locking protocols. Deadlocks are a necessary evil.
- **Starvation** is also possible if concurrency control manager is badly designed. For example:
 - A transaction may be waiting for an X-lock on an item, while a sequence of other transactions request and are granted an S-lock on the same item.
 - The same transaction is repeatedly rolled back due to deadlocks.

T1	T2	T3	T4
	lock-S(A)		
lock-X(A)	...	lock-S(A)	
Wait...	lock-S(A)
Wait...

The Two-Phase Locking Protocol

- This is a protocol which ensures conflict-serializable schedules.
- Phase 1: Growing Phase
 - transaction may obtain locks
 - transaction may not release locks
- Phase 2: Shrinking Phase
 - transaction may release locks
 - transaction may not obtain locks
- The protocol assures serializability. It can be proved that the transactions can be serialized in the order of their **lock points** (i.e. the point where a transaction acquired its final lock).

The Two-Phase Locking Protocol (Cont.)

- Two-phase locking *does not* ensure freedom from deadlocks

T_3	T_4
lock-X(B) read(B) $B := B - 50$ write(B) lock-X(A)	lock-S(A) read(A) lock-S(B)

Partial Schedule Under Two-Phase Locking

- Cascading roll-back is possible under two-phase locking.
- To avoid this, follow a modified protocol called **strict two-phase locking**. Here a transaction must hold all its **exclusive** locks till it commits/aborts.
- **Rigorous two-phase locking** is even stricter: here **all** locks are held till commit/abort. In this protocol transactions can be serialized in the order in which they commit.

T_5	T_6	T_7
lock-X(A) read(A) lock-S(B) read(B) write(A) unlock(A)		
	lock-X(A) read(A) write(A) unlock(A)	lock-S(A) read(A)

事务 T_7 的 read(A)步骤之后事务 T_5 发生故障，从而导致 T_6 与 T_7 级联回滚。

Lock Conversions

■ Two-phase locking with lock conversions:

- First Phase:

- can acquire a lock-S on item
- can acquire a lock-X on item
- can convert a lock-S to a lock-X (upgrade)

- Second Phase:

- can release a lock-S
- can release a lock-X
- can convert a lock-X to a lock-S (downgrade)

■ This protocol assures serializability. But still relies on the programmer to insert the various locking instructions.

Incomplete Schedule With a Lock Conversion

- T8:

- Read(a1)
- Read(a2)
-
- Read(an)
- Write(a1)

- T9:

- Read(a1)
- Read(a2)
- Display(a1 + a2)

T_8	T_9
lock-S(a_1)	
lock-S(a_2)	lock-S(a_1)
lock-S(a_3)	lock-S(a_2)
lock-S(a_4)	
	unlock(a_1)
	unlock(a_2)
lock-S(a_n)	
upgrade(a_1)	


T8 and T9 can access a1 and a2 simultaneously.

- If we employ the two-phase locking protocol, then T8 must lock a1 in X-mode. Then any concurrent execution of both transactions amounts to a serial execution.

Implementation of Locking

- A **lock manager** can be implemented as a separate process to which transactions send lock and unlock requests.
- The lock manager replies to a lock request by sending a lock grant messages (or a message asking the transaction to roll back, in case of a deadlock).
- The requesting transaction waits until its request is answered.
- The lock manager maintains a data-structure called a **lock table** to record granted locks and pending requests.
- The lock table is usually implemented as an in-memory hash table indexed on the name of the data item being locked.

Lock Table



- Deep blue rectangles indicate granted locks, light ones indicate waiting requests
- Lock table also records the type of lock granted or requested
- New request is added to the end of the queue of requests for the data item, and granted if it is compatible with all earlier locks
- Unlock requests result in the request being deleted, and later requests are checked to see if they can now be granted
- If transaction aborts, all waiting or granted requests of the transaction are deleted
 - lock manager may keep a list of locks held by each transaction, to implement this efficiently

15.2 Deadlock Handling

- A system is in a **deadlock state** if there exists a set of transactions such that every transaction in the set is waiting for another transaction in the set.
- Schedule with deadlock

T_1	T_2
lock-X on A write (A) wait for lock-X on B	lock-X on B write (B) wait for lock-X on A

- There are two principal methods:
 - **deadlock prevention** protocol to ensure that the system will never enter a deadlock state
 - allow the system to enter a deadlock state, and then try to recover by using a **deadlock detection** and **deadlock recovery** scheme

15.2.1 Deadlock Handling

■ **Deadlock prevention** protocols ensure that the system will never enter into a deadlock state. Some prevention strategies :

- 要求每个事务在开始之前封锁它的所有数据项(先声明)
 - ▶ 事务开始之前很难预知所需要的所有数据项
 - ▶ 导致数据项的使用率很低
- 对所有数据项强加一个次序，事务只能按要求的次序顺序封锁数据项
 - ▶ 例如：两阶段锁协议

More Deadlock Prevention Strategies

- The second approach for preventing deadlocks is to use preemption (抢占) and transaction rollbacks (事务回滚).
 - **wait-die** scheme — non-preemptive
 - ▶ older transaction may wait for younger one to release data item. Younger transactions never wait for older ones; they are rolled back instead.
 - ▶ a transaction may die several times before acquiring needed data item
 - **wound-wait** scheme — preemptive
 - ▶ older transaction *wounds* (forces rollback) of younger transaction instead of waiting for it. Younger transactions may wait for older ones.
 - ▶ may be fewer rollbacks than *wait-die* scheme.
- Both in wait-die and in wound-wait schemes, a rolled back transaction is restarted with its original timestamp. Older transactions thus have precedence over newer ones, and starvation is hence avoided.


Deadlock prevention (Cont.)

- Another simple approach to deadlock prevention is based on **lock timeouts**:
 - a transaction waits for a lock only for a specified amount of time. After that, the wait times out and the transaction is rolled back.
 - thus deadlocks are not possible
 - simple to implement; but starvation is possible. Also difficult to determine good value of the timeout interval.


15.2.2 Deadlock Detection

- Deadlocks can be described as a *wait-for graph*, which consists of a pair $G = (V, E)$,
 - V is a set of vertices (all the transactions in the system)
 - E is a set of edges; each element is an ordered pair $T_i \rightarrow T_j$.
- If $T_i \rightarrow T_j$ is in E , then there is a directed edge from T_i to T_j , implying that T_i is waiting for T_j to release a data item.
- When T_i requests a data item currently being held by T_j , then the edge $T_i \rightarrow T_j$ is inserted in the wait-for graph. This edge is removed only when T_j is no longer holding a data item needed by T_i .
- The system is in a deadlock state if and only if the wait-for graph has a cycle. Must invoke a deadlock-detection algorithm periodically to look for cycles.

Deadlock Detection (Cont.)



Wait-for graph without a cycle



Wait-for graph with a cycle

Deadlock Recovery

■ When deadlock is detected :

- **Selection of a victim.** Some transaction will have to rolled back (made a victim) to break deadlock. Select that transaction as victim that will incur minimum cost.
 - ▶ a. How long the transaction has computed, and how much longer the transaction will compute before it completes its designated task.
 - ▶ b. How many data items the transaction has used.
 - ▶ c. How many more data items the transaction needs for it to complete.
 - ▶ d. How many transactions will be involved in the rollback.
- **Rollback** -- determine how far to roll back transaction
 - ▶ Total rollback: Abort the transaction and then restart it.
 - ▶ More effective to roll back transaction only as far as necessary to break deadlock.
- **Starvation** happens if same transaction is always chosen as victim. Include the number of rollbacks in the cost factor to avoid starvation

15.4 Timestamp-Based Protocols

- Each transaction is issued a timestamp when it enters the system. If an old transaction T_i has time-stamp $\text{TS}(T_i)$, a new transaction T_j is assigned time-stamp $\text{TS}(T_j)$ such that $\text{TS}(T_i) < \text{TS}(T_j)$.
- The protocol manages concurrent execution such that the time-stamps determine the serializability order.
- In order to assure such behavior, the protocol maintains for each data Q two timestamp values:
 - **W-timestamp(Q)** 表示成功执行**write(Q)**的所有事务的最大时间戳.
 - **R-timestamp(Q)**表示成功执行**read(Q)**的所有事务的最大时间戳.

Timestamp-Based Protocols (Cont.)

- 基于时间戳的排序协议可以保证任何有冲突的**read** 和 **write** 操作按时间戳顺序执行。
- Suppose a transaction T_i issues a **read(Q)**:
 1. If $TS(T_i) \leq W\text{-timestamp}(Q)$, then T_i needs to read a value of Q that was already overwritten.
 - ▶ Hence, the **read** operation is rejected, and T_i is rolled back.
 2. If $TS(T_i) \geq W\text{-timestamp}(Q)$, then the **read** operation is executed, and $R\text{-timestamp}(Q)$ is set to **max(R-timestamp(Q), TS(T_i))**.

Timestamp-Based Protocols (Cont.)

■ Suppose that transaction T_i issues **write(Q)**.

1. If $\text{TS}(T_i) < \text{R-timestamp}(Q)$, then the value of Q that T_i is producing was needed previously, and the system assumed that that value would never be produced.
 - ▶ Hence, the **write** operation is rejected, and T_i is rolled back.
2. If $\text{TS}(T_i) < \text{W-timestamp}(Q)$, then T_i is attempting to write an obsolete value of Q.
 - ▶ Hence, this **write** operation is rejected, and T_i is rolled back.
3. Otherwise, the **write** operation is executed, and $\text{W-timestamp}(Q)$ is set to $\text{TS}(T_i)$.


Example Use of the Protocol

A partial schedule for several data items for transactions with timestamps 1, 2, 3, 4, 5. (Analysis it by yourself)

T_1	T_2	T_3	T_4	T_5
				read (X)
read (Y)	read (Y)	write (Y) write (Z)		read (Z)
read (X)	read (Z) abort		read (W)	
		write (W) abort		write (Y) write (Z)

Correctness of Timestamp-Ordering Protocol

- The timestamp-ordering protocol guarantees serializability since all the arcs in the precedence graph are of the form:



Thus, there will be no cycles in the precedence graph.

- Timestamp protocol ensures freedom from deadlock as no transaction ever waits.
- But the schedule may not be cascade-free, and may not even be recoverable.

End of Chapter

Assignment

- 15.21,29
- A scheduled flight has 50 tickets. Agency one wants to book 10 tickets and agency two want to book 20 tickets. Please design a lock strategy to implement the concurrency control.

Agency one	Agency two
Read(A)	
	Read(A)
A=A-10	
Write(A)	
	A=A-20
	Write(A)